

**CSCI 6333/6315 Database
Final Exam, Spring 2020**

Name: _____

ID: _____

Note: Please type your answers. If you need to draw a figure, you may draw on a paper, take a picture, and then cut and paste the figure to your file.

Problem 1 [20 points]. Consider the schema $R = (A, B, C, D, E, F)$ and the following set F of functional dependencies holds on R :

$$\begin{aligned}A &\rightarrow BCD \\ BC &\rightarrow DE \\ B &\rightarrow D \\ D &\rightarrow A\end{aligned}$$

- a) Compute A^+ .

Answer:

$$\begin{aligned}A &\Rightarrow A, A \rightarrow A \\ &\Rightarrow ABCD, A \rightarrow BCD \\ &\Rightarrow ABCDE, BC \rightarrow DE\end{aligned}$$

Hence, $A^+ = ABCDE$.

- b) Prove that the decomposition $R_1 = (A, B, C, F)$ and $R_2 = (B, D, E, F)$ of R is a lossless-join decomposition.

Answer:

$$\begin{aligned}R_1 \cap R_2 &= BF \\ BF &\Rightarrow BF, BF \rightarrow BF \\ &\Rightarrow BDF, B \rightarrow D \\ &\Rightarrow ABDF, D \rightarrow A \\ &\Rightarrow ABCDF, A \rightarrow BCD \\ &\Rightarrow ABCDEF, BC \Rightarrow DE\end{aligned}$$

Hence, $(BF)^+ = ABCDEF$, which implies $BF \rightarrow R_1$ and $BF \rightarrow R_2$, therefore the decomposition is lossless.

Problem 2. [20 points] Let the relations $r_1(A, B, C)$ and $r_2(C, D, E)$ have the following properties: r_1 has 30,000 tuples, r_2 has 45,000 tuples, 25 tuples of r_1 fit on one block, and 30 tuples of r_2 fit on one block. Estimate the number of block accesses required, using each of the following join strategies for $r_1 \bowtie r_2$ to determine the efficient loop

structure, that is, which relation shall be used for the outer loop (and the other is for the inner loop).

- a) Nested-loop join
- b) Block nested-loop join

Solution:

$$n_{r1} = 30000, \quad n_{r2} = 45000, \quad b_{r1} = \frac{n_{r1}}{25} = 1200, \quad b_{r2} = \frac{n_{r2}}{30} = 1500$$

- a) When r1 is for the outer and r2 is for the inner loop, the estimated cost is $n_{r1} * b_{r2} + b_{r1} = 30000 * 1500 + 1200 = 45,001,200$ block transfers, plus $n_{r1} + b_{r1} = 30,000 + 1,200 = 31,200$ seeks.

When r2 is for the outer loop and r1 is for the inner loop, the estimated cost is $n_{r2} * b_{r1} + b_{r2} = 45,000 * 1,200 + 1,500 = 54,001,500$ block transfers, plus $n_{r2} + b_{r2} = 45,000 + 1,500 = 46,500$ seeks.

It is more efficient to have r1 for the outer loop and for r2 for the inner loop.

- b) When r1 is for the outer and r2 is for the inner loop, the estimated cost is $b_{r1} * b_{r2} + b_{r1} = 1,200 * 1,500 + 1,200 = 1,801,200$ block transfers, plus $2 * b_{r1} = 2 * 1,200 = 2,400$ seeks.

When r2 is for the outer loop and r1 is for the inner loop, the estimated cost is $b_{r2} * b_{r1} + b_{r2} = 1,500 * 1,200 + 1,500 = 1,801,500$ block transfers, plus $2 * b_{r2} = 2 * 1,500 = 3,000$ seeks.

It is more efficient to have r1 for the outer loop and for r2 for the inner loop.

Problem 3. [10 points] Give two concrete examples to show that there are schedules that are possible under the two-phase locking protocol, but are not possible under the timestamp protocol, and *vice versa*.

Answer:

- a) A 2PL protocol but not timestamp

T3	T4
Lock(A)	
W(A)	
	Lock(B)
	W(B)
	Unlock(B)

Lock(B)
R(B)
Unlock(B)
Unlock(A)

b) A timestamp protocol but not 2PL

T5	T6
W(A)	
W(B)	<p>← for 2PL, T5 must unlock B so that T6 lock B and then write B</p>
	W(B)
	<p>← for 2PL, T5 must unlock A so that T6 can lock A and then read A</p>
	R(A)
	<p>← for 2PL, T6 must unlock A so that T5 can lock A and then read A</p> <p>This contradicts to 2PL for T5</p>
R(A)	

Problem 4. [10 points] When the system recovers from a crash, it constructs an undo-list and a redo-list. Explain why log records for transactions on the undo list must be processed in reverse order, while those log records for transactions on the redo-list are processed in a forward direction.

Answer:

Consider the following log records:

<T1, x, 10, 20>
 <T1, x, 20, 30>
 <T1, x, 30, 40>

To undo, we'd like to restore the oldest value 10 to x. If we do so in forward order, the value of x will be 30. If we do so in backward order, x will be correctly set to 10.

To redo, if we follow the backward order, we'll set x to 20, which shall be 40. If we follow the forward order, we'll correctly set x to 40.

Problem 5. [10 points] Let R be a relation and F a set of functional dependencies of R . $|R| = n$ and $|F| = m$.

- a) What is the time complexity to check whether R is in BCNF?
- b) What is the time complexity to check whether R is in 3NF?

Answer:

- a) Following the definition of BCNF, we need to check, for any nontrivial functional dependency $\alpha \rightarrow \beta$ in F , whether $\alpha \rightarrow R$, i.e., whether α is superkey for R . To do so, we need to compute α^+ and to check whether it contains R . We need in the worst case to compute one attribute closure for every functional dependency in F . Since the time complexity of computing an attribute closure is $O(mn^2)$, so the total time needed is $O(m^2n^2)$.
- b) Unlike the problem of testing whether R is BCNF, the problem of testing R is in 3NF is hard. For any nontrivial functional dependency $\alpha \rightarrow \beta$ where α is not a superkey for R , the third condition of the 3NF definition asks for each attribute A in $\beta - \alpha$ whether A is contained in a candidate key for R . This implies that to facilitate the testing for the third conditions, we need to find candidate keys for R , which is NP-hard. Hence, the time complexity of testing whether R is in 3NF is NP-hard.

Problem 6. [10 points] A disk block has 1024 bytes, and both a pointer and a search key are of 4 bytes each. Now, consider that we build a B^+ -tree index with a disk block to store every tree node. Assume that the index tree has a height of 4. Estimate the number of unique search key values that can be represented by the index tree.

Answer:

Following the given information, a block can store $1024/4 = 256$ pointer and/or search key values. Thus, with the size of a disk block, each internal node can store 128 pointers and 127 search key values, leaving four bytes wasted.

Each internal node has at least $\lceil 128/2 \rceil = 64$ children and at most 128 children. With height 4, the B^+ -tree index has at least 64^4 leaf nodes and at most 128^4 leaf nodes.

At leaf a node, the 4 bytes wasted in an internal node can be used to point the next leaf node. By this way, each leaf node has at least $\lceil 128/2 \rceil = 64$ search key values and at most 128^4 search key values.

Therefore, with height 4, the B^+ -tree index will have at least $64^5 = 1,073,741,824$ search key values and at most $128^5 = 34,359,738,368$ search key values.

Problem 7 [20 points] Show that the wound-wait strategy will prevent

- a) Deadlock, and
- b) Starvation.

Proof.

Let T_1, T_2, \dots, T_n be a list of transactions. Suppose that they follow the wound-wait strategy with timestamp $t(T_i) = i$. We have the following two claims.

Claim 1: There is no deadlock.

Proof of Claim 1. We construct a wait-for graph G for T_1, T_2, \dots, T_n . Assume by contradiction that there is a deadlock. That is, G has a cycle $T_{i_1} \rightarrow T_{i_2} \rightarrow \dots \rightarrow T_{i_j} \rightarrow T_{i_1}$. Because all transactions follow the wound-wait strategy, and this strategy only allows a younger transaction to wait for an older transaction, the cycle implies

$$t(T_{i_1}) > t(T_{i_2}) > \dots > t(T_{i_j}) > t(T_{i_1}).$$

Thus, we have a contradiction. Hence, there is no deadlock.

Claim 2: There is no starvation.

Proof of Claim 2. Let $f(T_i)$ denote the time at which transaction T_i is completely executed. Let us consider T_1 first. Note that T_1 is the oldest transaction. When T_1 became active, there are two possibilities: (a) There were no other active transactions; and (b) There were some active transactions. For (a), since T_1 is the oldest transaction, it never be wounded by any other transaction, so it will continue execution till completion. Thus, it will not starve and will finish at time $f(T_1)$. For (b), since T_1 is the oldest, by the wound-die strategy it will wound all active transactions to force them to abort so as to gain its execution. Thus, this enters case (a). Hence, T_1 will not starve and finish at time $f(T_1)$.

For transaction T_2 , it may be completed before $f(T_1)$. If not, then after $f(T_1)$, T_2 becomes the oldest transaction and, by the similar analysis for T_1 , T_2 it will not starve and complete its execution at time $f(T_2)$. We can carry out the similar analysis for T_3, \dots, T_n , and none of them will starve.