

**CSCI 6315 Applied Database Systems**  
**ASSIGNMENT 5: Transaction Management**  
**Due is 05/03/2020 00:00**  
**Ulvi Bajarani**  
**Student ID 20539914**  
**E-mail: [ulvi.bajarani01@utrgv.edu](mailto:ulvi.bajarani01@utrgv.edu)**

## Questions and Answers:

**Problem 1.** Explain the distinction between the terms serial schedule and serializable schedule.

**Answer 1.** While serial schedule contains the transactions that complete one after another (for example,  $T_1$ , then  $T_2$ , then  $T_3$ , etc.), the serializable schedule is the *non-serial schedule* (that contains transactions loading concurrently: this might lead to some problems.) which, however, might be transformed to the serial schedule. It should be mentioned that not all non-serial schedules are serializable.

**Problem 2.** Consider the following two transactions:

$T_1$ :	read( $A$ ); read( $B$ ); if $A = 0$ then $B = B + 1$ ; write $B$ ;
$T_2$ :	read( $B$ ); read( $A$ ); if $B = 0$ then $A = A + 1$ ; write $A$ ;

Figure 1. Two transactions  $T_1$  and  $T_2$

Let the consistency requirement be  $A = 0 \vee B = 0$  with  $A = B = 0$  the initial values.

- a. Show that every serial execution involving these two transactions preserves the consistency of the database.
- b. Show a concurrent execution of  $T_1$  and  $T_2$  that produces a nonserializable schedule.
- c. Is there a concurrent execution of  $T_1$  and  $T_2$  that produces a serializable schedule?

**Answer 2.**

- a. While  $T_1$  executes first, the if branch in  $T_2$  returns false, and vice versa. The consistency requirement is  $A = 0 \vee B = 0$  with  $A = B = 0$  the initial values, so one of either  $A$  or  $B$  remains 0, which allows to preserve the consistency.
- b. The table below might be the solution. Swapping of the order of read( $B$ ); and write  $B$ ; leads the different results :

$T_1$	$T_2$
read( $A$ );	
	read( $B$ );
	read( $A$ );
read( $B$ );	
if $A = 0$ then $B = B + 1$ ;	
	if $B = 0$ then $A = A + 1$ ;
write $B$ ;	
	write $A$ ;

- c. There is no solution, because in the concurrent case, writes will not be performed. As a result, the reads in either  $T_2$  and  $T_1$  gets the old values, not changed

**Problem 3.** Consider the precedence graph in Figure 2. Is the corresponding schedule serializable?




Figure 2. Precedence graph

**Answer 3.** The corresponding schedule is serializable, because there is no cycle between the transactions. The serial schedule is the topological sorting, for example,  $T_1, T_2, T_3, T_4, T_5$ .

**Problem 4.** Show that two-phase locking protocol ensures conflict serializability, and that transaction can be serialized according to their lock points.

**Answer 4.** If 2PL didn't ensure the conflict serializability, there would be  $T_0, T_1, T_2, T_3, T_4 \dots T_n$  transactions that are non-serializable. In that case, there should be the cycle in the precedence graph, which should lead to  $T_0, T_1, T_2, \dots T_1, T_0$  precedence graph. Let  $\alpha_i$  be the time of occurred lock. If 2PL existed in the cyclic precedence graph, the locking time should be  $\alpha_0 < \alpha_1 \dots < \alpha_0$ , which is the contradiction. Thus, 2PL cannot create non-serializable schedule. Also,  $T_i \rightarrow T_j$  and  $\alpha_i < \alpha_j$  means that the lock point ordering is a topological sort ordering the precedence graph, which means the serializability.

**Problem 5.** Consider transactions  $T_1$  and  $T_2$  in Figure 1. Add lock and unlock instructions to them so that they observe the two-phase locking protocol. Can the execution of these two transactions result in a deadlock?

**Answer 5.** The 2PL-protocol-observing case:

$T_1$	$T_2$
Lock-S( $A$ );	
read( $A$ );	
Lock-X( $B$ );	
read( $B$ );	
if $A = 0$ then $B = B + 1$ ;	
write $B$ ;	
Unlock-S( $A$ );	
Unlock-X( $B$ );	
	Lock-S( $B$ );
	read( $B$ );
	Lock-X( $A$ );
	read( $A$ );
	if $B = 0$ then $A = A + 1$ ;
	write $A$ ;
	Unlock-S( $B$ );
	Unlock-X( $A$ );

The deadlock case. Neither  $T_1$  nor  $T_2$  might make a progress:

$T_1$	$T_2$
Lock-S( $A$ );	
read( $A$ );	
	Lock-S( $B$ );
	read( $B$ );
Lock-X( $B$ );	
read( $B$ );	
	Lock-X( $A$ );
	read( $A$ );
if $A = 0$ then $B = B + 1$ ;	
write $B$ ;	
Unlock-X( $B$ );	
Unlock-S( $A$ );	
	if $B = 0$ then $A = A + 1$ ;
	write $A$ ;
	Unlock-S( $B$ );
	Unlock-X( $A$ );

**Problem 6.** Show that there are schedules that are possible under the two-phase locking protocol, but are not possible under the timestamp protocol, and vice versa.

**Answer 6.** the schedule with 2PL that is not possible in the Timestamp Protocol (due to step 5, where  $TS(T_0) < W\text{-Timestamp}(B)$  ):

$T_0$	$T_1$
Lock-S( $A$ );	Lock-X( $B$ );
read( $A$ );	write $B$ ;
	unlock( $B$ );
Lock-S( $B$ );	
read( $B$ );	
unlock( $A$ );	
unlock( $B$ );	

the schedule with Timestamp Protocol that is not possible in the 2PL Protocol, because it requires  $T_1$  to unlock  $A$  between the steps 2 and 3, and lock  $B$  between steps 4 and 5:

$T_0$	$T_1$	$T_2$
write $A$ ;		write $A$ ;
	write $A$ ;	
write $B$ ;		
	write $B$ ;	

**Problem 7.** Explain the purpose of the checkpoint mechanism. How often should checkpoints be performed? How does the frequency of checkpoints affect

- System performance when no failure occurs
- The time it takes to recover from a system crash



- The time it takes to recover from a disk crash

**Answer 7.**

- Decreases the system performance (negatively);
- Decreases the time of system recovery (positively);
- Decreases the time of disk recovery (positively);

**Problem 8.** When the system recovers from a crash, it constructs an undo-list and a redolist. Explain why log records for transactions on the undo list must be processed in reverse order, while those log records for transactions on the redo-list are processed in a forward direction.

**Answer 8.** The answer is trivial: if the undo-list was processed in the forward order, it would give wrong values to the data that is updated with several transactions. The same is right for the redo-list processed in the reverse direction.